

Introduction to Operating Systems

Chapter 3: Interprocess communication

Manuel

Fall 2017

Outline

- 1 Exhibiting the problem
- 2 Solving the problem
- 3 More solutions

Easy setup

Independent threads:

- Cannot affect/be affected by anything
- State not shared with other threads
- Input state determines the output
- Reproducible
- No side effect when stopping/resuming

Executing threads

Where difficulty starts:

- Single-tasking: run a thread to completion and start next one
- Multi-tasking:
 - One core shared among several threads
 - Several cores run several threads in parallel
- A thread runs on one core at a time
- A thread can run on different cores at different times
- For threads no difference between one or more cores

Cooperating threads

Setup:

- Threads sharing a state
- Behavior depends on the execution sequence
- Behavior may seem random/irreproducible

Cooperating threads

Setup:

- Threads sharing a state
- Behavior depends on the execution sequence
- Behavior may seem random/irreproducible

Major problems:

- ① How can threads/processes share information?
- ② How to prevent them for getting in each other's way?
- ③ How to handle sequencing?


Atomic operation

Operation that either happens in its entirety or not at all

- Atomic operation cannot be created
- Atomic operations are hardware level
- Central to the question of parallel programming

Example of atomic operation: $A=B$, read a clean value for B and sets a clean value for A

Race conditions



- Process A wants to queue a file, reads `next_free_slot=7`
- Interrupt occurs, Process B reads `next_free_slot=7`
- Process B queues its file in slot 7, and update `next_free_slot=8`
- Process A resumes using `next_free_slot=7`
- Game over

Too much milk



9AM

9:15AM

9:30AM

9:45AM

Too much milk




9AM

9:15AM

9:30AM

9:45AM

Too much milk



9AM

9:15AM

9:30AM

9:45AM

Too much milk



9AM

9:15AM

9:30AM

9:45AM

Too much milk



9AM

9:15AM

9:30AM

9:45AM

Too much milk



9AM

9:15AM

9:30AM

9:45AM

Outline

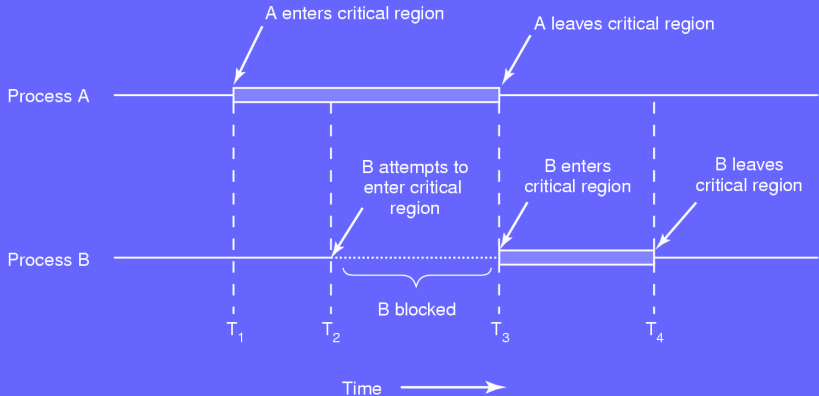
- 1 Exhibiting the problem
- 2 Solving the problem
- 3 More solutions

Critical region

Part of the program where shared memory is accessed:

- No two processes can be in a critical region at the same time
- No assumption on speed/number of CPU
- No process outside a critical region can block other processes
- No process waits forever to enter its critical region

Mutual exclusion



Too much milk

Frank

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

John

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

Too much milk

Frank

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

John

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

Result: too much milk

Too much milk

Frank

```
1 leave note Frank;  
2 if(no note John) {  
3     if(no milk){  
4         milk the cow;  
5     }  
6 }  
7 remove note Frank;
```

John

```
1 leave note John;  
2 if(no note Frank) {  
3     if(no milk) {  
4         milk the cow;  
5     }  
6 }  
7 remove note John;
```

Too much milk

Frank

```
1 leave note Frank;  
2 if(no note John) {  
3     if(no milk){  
4         milk the cow;  
5     }  
6 }  
7 remove note Frank;
```

John

```
1 leave note John;  
2 if(no note Frank) {  
3     if(no milk) {  
4         milk the cow;  
5     }  
6 }  
7 remove note John;
```

Result: no milk

Too much milk

Frank

```
1 leave note Frank;
2 while(note John) {
3     nothing;
4 }
5 if(no milk) {
6     milk the cow;
7 }
8 remove note Frank;
```

John

```
1 leave note John;
2 if(no note Frank) {
3     if(no milk) {
4         milk the cow;
5     }
6 }
7 remove note John;
```

Too much milk

Frank

```
1 leave note Frank;  
2 while(note John) {  
3     nothing;  
4 }  
5 if(no milk) {  
6     milk the cow;  
7 }  
8 remove note Frank;
```

John

```
1 leave note John;  
2 if(no note Frank) {  
3     if(no milk) {  
4         milk the cow;  
5     }  
6 }  
7 remove note John;
```

Result: just enough milk

Too much milk

Comments on the solution:

- Solution is correct
- Complicated
- Asymmetrical
- Busy wait, CPU time wasted

Peterson's idea

Simple strategy:

- Assume two processes
- Two functions:
 - `enter_region`: show process is interested and wait for its turn
 - `leave_region`: indicates departure from critical region
- Solution is correct
- Drawback: busy wait

Peterson's idea

Simple strategy:

- Assume two processes
- Two functions:
 - `enter_region`: show process is interested and wait for its turn
 - `leave_region`: indicates departure from critical region
- Solution is correct
- Drawback: busy wait

Exercise: “implement” Peterson's idea (use two processes represented by the integers 0 and 1)

Peterson's idea

```
1  #define TRUE 1
2  #define FALSE 0
3  int turn;
4  int interested[2];
5  void enter_region(int p) {
6      int other;
7      other=1-p;
8      interested[p]=TRUE;
9      turn=p;
10     while(turn==p && interested[other]==TRUE)
11 }
12 void leave_region(int p) {
13     interested[p]=FALSE;
14 }
```

Peterson's idea

Side effects of Peterson's idea:

- Two processes: L, low priority, and H, high priority
- L enters in a critical region
- H becomes ready
- H has higher priority so the scheduler switches to H
- L has lower priority so is not rescheduled as long as H is busy
- H loops forever

Software vs. hardware?

Disabling interrupts, good or bad?

Software vs. hardware?

Disabling interrupts, good or bad?

- Case 1: block interrupts for the whole computer
 - Interrupts could be disabled for a while (too long?)
 - This gives a lot of power to the programmer
- Case 2: block interrupts for only one CPU
 - No effect on other processors
 - Another CPU can access the variable between read and write

The TSL instruction

- Test and Set Lock (TSL)
- Hardware level, requires assembly
- Task: copy *lock* to register and set *lock* to 1
- Atomic operation

The TSL instruction

- Test and Set Lock (TSL)
- Hardware level, requires assembly
- Task: copy *lock* to register and set *lock* to 1
- Atomic operation

```
1  enter_region:
2      TSL REGISTER, LOCK
3      CMP REGISTER, #0
4      JNE enter_region
5      RET
6
7  leave_region:
8      MOVE LOCK, #0
9      RET
```


The TSL instruction

- Test and Set Lock (TSL)
- Hardware level, requires assembly
- Task: copy *lock* to register and set *lock* to 1
- Atomic operation

```
1  enter_region:
2      TSL REGISTER, LOCK
3      CMP REGISTER, #0
4      JNE enter_region
5      RET
6
7  leave_region:
8      MOVE LOCK, #0
9      RET
```

Note: TSL displays the same side-effect as Peterson's solution when dealing with processes having different priorities

Blocking strategy

```
1  #define N 100
2  int count=0;
3  void producer() {
4      int item;
5      while(1) {
6          item=produce_item();
7          if(count==N) sleep();
8          insert_item(item); count++;
9          if(count==1) wakeup(consumer);
10     }
11 }
12 void consumer() {
13     int item;
14     while(1) {
15         if(count==0) sleep();
16         item=remove_item(); count--;
17         if(count==N-1) wakeup(producer);
18         consume_item(item);
19     }
20 }
```

Race condition

Setup: buffer is empty

- Consumer reads *count* == 0
- Scheduler stops the consumer and starts the producer
- Producer adds one item
- Producer wakes up the consumer
- Consumer not yet asleep, signal is lost
- Consumer goes asleep
- When the buffer is full the producer falls asleep
- Both consumer and producer sleep forever

Semaphore

Basics:

- 1965, Edsger Dijkstra
- Simple hardware based solution
- Basis of all contemporary OS synchronization mechanisms

Semaphore

Basics:

- 1965, Edsger Dijkstra
- Simple hardware based solution
- Basis of all contemporary OS synchronization mechanisms

A semaphore s is a non-negative variable that can only be changed or tested by two atomic actions:

```
1 down(s) {  
2     while(s==0) wait();  
3     s--;  
4 }
```

```
1 up(s) {  
2     s++;  
3 }
```

Semaphore and processes

The down operation

Check if the value is > 0

- True: decrement the value and continues
- False: sleep, do not complete the down

Only one single atomic action

The up operation

- Increment the value of the semaphore
- If one or more processes were asleep, randomly choose one to wakeup (complete its down)

Only one single atomic action

Semaphore implementation

Semaphores MUST be implemented in an indivisible way

- Up and down are implemented using system calls
- OS disables all interrupts while testing, updating the semaphore and putting process to sleep
- When dealing with more than one CPU, semaphores are protected using the TSL instruction

Note: a semaphore operation only takes a few microseconds

Semaphore and interrupts

Hiding interrupts using semaphores:

- Each I/O device gets a semaphore initialised to 0
- Managing process applies a `down` when starting an I/O device
- Process is blocked
- Interrupt handler applies an `up` when receiving an interrupt
- Process is ready to run again

Mutex

MUTual EXclusion:

- Simplified semaphore
- Takes values 0 (unlocked) or 1 (locked)

On a mutex-lock request:

- If mutex is currently unlocked, lock it; thread can enter in critical region
- If mutex is currently locked, block the calling thread until thread in critical regions is done and calls mutex-unlock
- Randomly chose which thread acquires the lock if more than one are waiting

Mutex implementation

Mutexes can be implemented in user-space using TSL

```
1  mutex-lock:
2      TSL REGISTER,MUTEX
3      CMP REGISTER,#0
4      JNE ok
5      CALL thread_yield
6      JMP mutex-lock
7  ok: RET
8
9  mutex-unlock:
10     MOVE MUTEX,#0
11     RET
```

- What differences were introduced compared to `enter_region` (3.20)?
- In user-space what happens if a thread tries to acquire lock through busy-waiting?
- Why is `thread_yield` used?

Example

consumer-producer.c

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #define MAX 1000
4  pthread_mutex_t m; pthread_cond_t cc, cp; int buf=0;
5  void *prod(void *p) {
6      for(int i=1;i<MAX;i++) {
7          pthread_mutex_lock(&m);
8          while(buf!=0) pthread_cond_wait(&cp,&m);
9          buf=1; pthread_cond_signal(&cc);
10         pthread_mutex_unlock(&m);
11     } pthread_exit(0);
12 }
13 void *cons(void *p) {
14     for(int i=1;i<MAX;i++) {
15         pthread_mutex_lock(&m);
16         while(buf==0) pthread_cond_wait(&cc,&m);
17         buf=0; pthread_cond_signal(&cp);
18         pthread_mutex_unlock(&m);
19     } pthread_exit(0);
20 }
21 int main() {
22     pthread_t p, c;
23     pthread_mutex_init(&m,0);
24     pthread_cond_init(&cc,0); pthread_cond_init(&cp,0);
25     pthread_create(&c,0,cons,0); pthread_create(&p,0,prod,0);
26     pthread_join(p,0); pthread_join(c,0);
27     pthread_cond_destroy(&cc); pthread_cond_destroy(&cp);
28     pthread_mutex_destroy(&m);
29 }
```

Example

Alter the previous program such as:

- To display information on the consumer and producer
- To increase the buffers to 100
- To have two consumers and one producer. In this case also print which consumer is active.

Outline

- 1 Exhibiting the problem
- 2 Solving the problem
- 3 More solutions

Dangers...

```
1  mutex mut = 0; semaphore empty = 100; semaphore full = 0;
2  void producer() {
3      while(TRUE) {
4          item = produce_item();
5          mutex-lock(&mut);
6          down(&empty);
7          insert_item(item);
8          mutex-unlock(&mut);
9          up(&full);
10     }
11 }
12 void consumer() {
13     while(TRUE) {
14         down(&full);
15         mutex-lock(&mut);
16         item = remove_item();
17         mutex-unlock(&mut);
18         up(&empty); consume_item(item);
19     }
20 }
```

Questions

In the previous code:

- What is the behavior of the producer when the buffer is full?
- What about the consumer?
- What is the final result for this program?
- How to fix it?

Monitors

Semaphores and mutexes are convenient but sometimes risky

The concept of monitor:

- Monitors are higher level, cleaner and less risky solution
- A monitor is composed of
 - Procedures/functions
 - Structures
- Only one process can be active within a monitor at a time
- Monitors are a programming concept, compiler must know them
- Mutual exclusion handled by the compiler not the programmer

Monitors

Basics on monitors:

- Monitors are easy to use for mutual exclusion
- Offer a condition to block “properly” when the buffer is full
- A signal on the condition variable can wake up a blocked process
- Only one process can be active in the monitor
- As soon as the signal on the condition is sent, exit the monitor
- Other process can resume

Example

```
1  monitor ProducerConsumer {
2      condition full, empty;
3      int count;
4      void insert(item) {
5          if (count == N) wait(full);
6          insert_item(item);
7          count++;
8          if (count==1) signal(empty);
9      }
10     void remove() {
11         if (count==0) wait(empty);
12         removed = remove_item;
13         count--;
14         if (count==N-1) signal(full);
15         return(removed);
16     }
17     count:= 0;
18 }
```

```
1  void ProducerConsumer::producer() {
2      while (TRUE) {
3          item = produce_item();
4          ProducerConsumer.insert(item);
5      }
6  }
7  void ProducerConsumer::consumer() {
8      while (TRUE) {
9          item=ProducerConsumer.remove();
10         consume_item(item)
11     }
12 }
```

Message passing

Limitation of semaphore/monitor: processes need to share some part of the memory. Distributed systems over a network consist of multiple CPU each with its own private memory


Message passing:

- `send(destination,&message)`
- `receive(source,&message)`, can either block or exit if nothing is received

Potential issues:

- Message lost (sending/acknowledging reception)
- No possible confusion on process names
- Security (authentication, traffic)
- Performance

Barriers



Useful for programs/problems where all processes must finish before going to the next phase

Key points

- Why is thread communication essential?
- What is a critical region?
- Do software solutions exist?
- What is an atomic operation?
- What are the two best and most common solutions?

Thank you!